

# YU CAO

Email: yc3390@nyu.edu Phone: +1(917)575-4323

Github:<https://github.com/Yucaao42>, Webpage:<https://yucao42.github.io/>

## EDUCATION

---

**New York University, Courant Institute of Mathematical Sciences** Sep. 2018-Dec. 2021  
*Master in Computer Science*

**Beihang University** Aug. 2011-Jul. 2018  
*Master and Bachelor in Instrumentation Science and Engineering*

## PUBLICATION AND TECH REPORTS

- 
- Fuqiang Zhou, **Yu Cao**, Xinming Wang, Fast and Resource-efficient Hardware Implementation of Modified Line Segment Detector. *IEEE Transactions on Circuits and Systems for Video Technology*, Co-First-Author
  - **Yu Cao**, Fuqiang Zhou, Minimal Non-linear Camera Pose Estimation Method Using Lines for SLAM Applications. *IEEE Winter Conference on Applications of Computer Vision 2018*
  - Cheng Tan, Zhichao Li, Jian Zhang, **Yu Cao**, Sikai Qi, Zherui Liu, Yibo Zhu, Chuanxiong Guo Serving DNN models with multi-instance gpus: A case of the reconfigurable machine scheduling problem, *Arxiv 2021*
  - Jianwen Jiang, **Yu Cao**, Lin Song, Shiwei Zhang, Yunkai Li, Ziyao Xu, Qian Wu, Chuang Gan, Chi Zhang, Gang Yu, Human centric spatio-temporal action localization *ActivityNet Workshop on CVPR 2018*

## EXPERIENCE

---

**Research Engineer/Senior Software Engineer, Bloomberg AI** Feb. 2022- present

- Collaborate with colleagues on production systems and applications.
- Design, experiment, and evaluate algorithms as well as models.
- Research and develop pricing algorithm for fix-income market.

**Research Intern, ByteDance Applied Machine Learning Systems** May. 2021- Aug. 2021

- Researched into scheduling algorithms(greedy and Monte Carlo Tree Search) to allocate GPU resources for DNN inference jobs to save up to 20% GPUs with Multi-Instance-GPU(MIG) feature.
- Researched into scheduling problems in pipeline DNN model serving with MIG by fitting the allocation scheduling problem into continuous optimization framework using gradient descent.

**Applied Scientist Intern, AWS AI Recognition, Amazon** June. 2019- Aug. 2019

- Built a multi-language scene-text localization model based on MaskRCNN detection model.
- Researched into using visual and language semantic embeddings extracted from pretrained unsupervised Auto-encoder to improve detection performance.

**Algorithm Developer Intern, Face++, Megvii Inc.** Nov. 2017- Jul. 2018

- CVPR 2018 ActivityNet Challenge of Spatio-temporal Action Localization(1st place)
  - Combined I3D deep learning model with Non-local module to express Actor-Target-Relationship in videos.

**Teaching Assistant, Graduate Distributed Systems, New York University** Aug. 2019- Dec. 2019

- Helped students on assignments of replication/Raft protocols, map-reduce, fault-tolerant key-value service.

## PROJECT

---

**Resource Efficient Real-time Streaming Video Analytic System** Oct. 2019- May. 2021

- Identified the challenge in current real-time video analytic jobs' scheduler that the processing time of frames could change because of the changing input rate or the content dependent processing logic.
- Built a distributed system(8k lines of C++) with dynamic scheduler that monitors analytic jobs' latency and throughput online and makes automatic adjustment when performance changes are detected.
- Tested on real-world gaming and traffic analysis workloads using Deep Neural Networks, our system has up to 64% more throughput(within tight latency SLO) compared with state-of-the-art system.

## AWARDS

---

• McCracken Fellowship by NYU GSAS 2018-2021

## SKILLS

- 
- **Programming:** Python, C/C++, VHDL, Verilog, Go, Pytest Pytorch, Pandas, Linux
  - **Distributed System:** Kafka, KServe, Kubernetes, Argo, Hera, Grafana, Faust, AWS S3
  - **Machine Learning:** Deep learning, Random Forest, GBDT, XGBoost, Time Series
  - **Quantitative Finance:** Fixed Income Bond Pricing, Backtesting, Pricing Evaluation, Municipal Bonds, Illiquid Bond Pricing, Security Similarity Modeling